

## Spring 2024: CS5720 Neural Networks & Deep Learning - ICP-7

### Assignment-7

NAME: Bhavana Billa

STUDENT ID:700756590

Github Link:

<https://github.com/BillaBhavana7/neuralN>

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.optimizers import RMSprop, Adam
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import classification_report, confusion_matrix
import warnings
warnings.filterwarnings("ignore")
```

```
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 [=====] - 2s 0us/step

```
classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
```

```
y_train = y_train.reshape(-1,)
```

```
# Reshape converting 2D to 1D
y_test = y_test.reshape(-1,)
y_train = y_train.reshape(-1,)
```

```
# This code normalization
x_train = x_train / 255.0
x_test = x_test / 255.0
```

```
x_train.shape
```

```
(50000, 32, 32, 3)
```

```
from tensorflow.keras import layers, models
lenet = keras.models.Sequential([
    keras.layers.Conv2D(32, kernel_size=3, strides=1, activation='relu', input_shape=(32,32,3), padding='same'), #C1
    keras.layers.Conv2D(32, kernel_size=3, strides=1, activation='relu', padding='same'), #C2
    keras.layers.MaxPooling2D(pool_size=2, strides=2), #S1
    keras.layers.Dropout(0.25),
    keras.layers.Conv2D(64, kernel_size=3, strides=1, activation='relu', padding='same'), #C3
    keras.layers.Conv2D(64, kernel_size=3, strides=1, activation='relu', padding='same'), #C4
    keras.layers.MaxPooling2D(pool_size=2, strides=2), #S2
    keras.layers.Dropout(0.25),
    keras.layers.Conv2D(128, kernel_size=3, strides=1, activation='relu', padding='same'), #C5
    keras.layers.Conv2D(128, kernel_size=3, strides=1, activation='relu', padding='same'), #C6
    keras.layers.MaxPooling2D(pool_size=2, strides=2), #S3
    keras.layers.Dropout(0.25),
    keras.layers.Flatten(), #Flatten
    keras.layers.Dense(512, activation='relu'), #F1
    keras.layers.Dropout(0.5),
    keras.layers.Dense(10, activation='softmax') #Output Layer
])
```

```
lenet.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_2 (Dropout)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 512)	1049088

```
lenet.compile(optimizer='adam', loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
```

```
hist = lenet.fit(x_train, y_train, epochs=5, validation_data=(x_test, y_test), verbose=1)
```

Epoch 1/5

1563/1563 [=====] - 397s 254ms/step - loss: 1.2135 - accuracy: 0.5651 - val\_loss: 1.0869 - val\_accuracy: 0.6111

Epoch 2/5

1563/1563 [=====] - 406s 260ms/step - loss: 1.0362 - accuracy: 0.6328 - val\_loss: 0.8905 - val\_accuracy: 0.6853

Epoch 3/5

1563/1563 [=====] - 403s 258ms/step - loss: 0.9405 - accuracy: 0.6669 - val\_loss: 0.8178 - val\_accuracy: 0.7106

Epoch 4/5

1563/1563 [=====] - 395s 253ms/step - loss: 0.8742 - accuracy: 0.6951 - val\_loss: 0.8158 - val\_accuracy: 0.7195

Epoch 5/5

1563/1563 [=====] - 390s 250ms/step - loss: 0.8228 - accuracy: 0.7124 - val\_loss: 0.7578 - val\_accuracy: 0.7349

```

aug_data = keras.preprocessing.image.ImageDataGenerator(
    rotation_range=10,
    width_shift_range=0.1,
    height_shift_range=0.1,
    zoom_range=0.1,
    horizontal_flip=True,
    fill_mode='nearest')

aug_data.fit(x_train)

from tensorflow.keras import layers, models
lenet = keras.models.Sequential([
    keras.layers.Conv2D(32, kernel_size=3, activation='relu', input_shape=(32,32,3), padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(32, kernel_size=3, activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPooling2D(pool_size=2),
    keras.layers.Dropout(0.25),
    keras.layers.Conv2D(64, kernel_size=3, activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(64, kernel_size=3, activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPooling2D(pool_size=2),
    keras.layers.Dropout(0.25),
    keras.layers.Conv2D(128, kernel_size=3, activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(128, kernel_size=3, activation='relu', padding='same'),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPooling2D(pool_size=2),
    keras.layers.Dropout(0.25),
    keras.layers.Flatten(),

```

```

# summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title("Accuracy by LeNet on CIFAR-10 Data")
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()


# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Loss by LeNet on CIFAR-10 Data')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'])
plt.show()

```

```
|: <matplotlib.legend.Legend at 0x7f75101e8160>
```



```
|: <matplotlib.legend.Legend at 0x7f75101e8d30>
```



```

from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
y_predictions= lenet.predict(x_test)
y_predictions.reshape(-1,)
y_predictions= np.argmax(y_predictions, axis=1)

confusion_matrix(y_test, y_predictions)

```

```

13/313 [=====] - 22s 71ms/step
array([[612,  54,  29,  14,  23,   4,  24,  11, 167,  62],
       [ 19, 823,   0,   4,   1,   0,   9,   1,  43, 100],
       [114,  36, 173,  76, 174,  89, 199,  39,  42,  58],
       [ 47,  61,  27, 256,  51, 155, 228,  35,  29, 111],
       [ 44,  29,  29,  29, 399,  24, 279,  97,  18,  52],
       [ 24,  20,  23, 163,  70, 403, 108,  85,  14,  90],
       [ 10,  19,  13,  41,  19,  16, 813,   2,  10,  57],
       [ 27,  28,   3,  36,  88,  72,  63, 537,   8, 138],
       [109, 111,   2,   6,   5,   2,  14,   4, 695,  52],
       [ 37, 321,   0,   5,   1,   0,  17,   4,  54, 561]])


```

```

# confusion matrix and accuracy
from sklearn.metrics import confusion_matrix, accuracy_score
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions))
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()

```

Confusion matrix



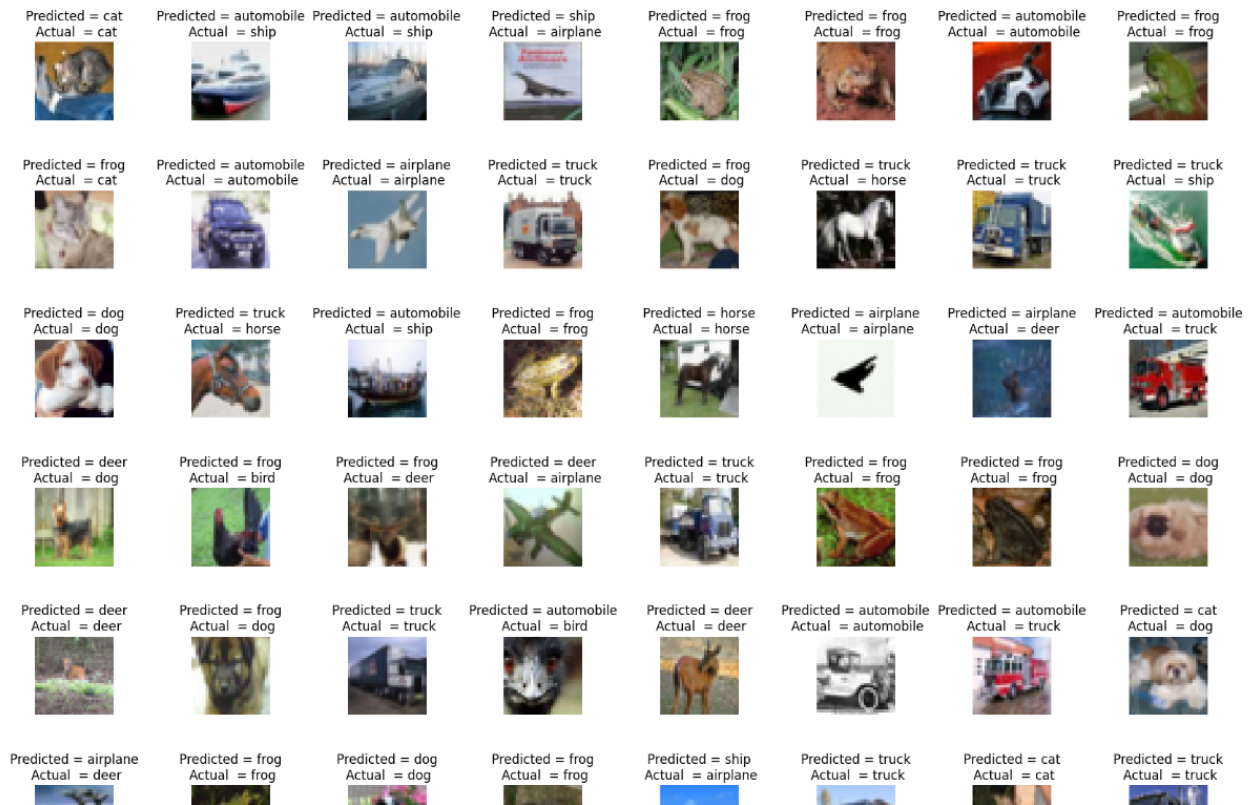
```
print("Test accuracy:", accuracy_score(y_test, y_predictions))
```

Test accuracy: 0.5272

```
L = 8
W = 8
fig, axes = plt.subplots(L, W, figsize = (20,20))
axes = axes.ravel() #

for i in np.arange(0, L * W):
    axes[i].imshow(x_test[i])
    axes[i].set_title("Predicted = {}\n Actual = {}".format(classes[y_predictions[i]], classes[y_test[i]]))
    axes[i].axis('off')

plt.subplots_adjust(wspace=1)
```



```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, Dropout, Flatten, MaxPooling2D
from tensorflow.keras.optimizers import SGD, Adam

```

```

#Define Alexnet Model
AlexNet = Sequential()
AlexNet.add(Conv2D(filters=16,kernel_size=(3,3),strides=(4,4),input_shape=(32,32,3), activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(5,5),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(30,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(20,(3,3),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Flatten())
AlexNet.add(Dense(200, activation='relu'))
AlexNet.add(Dropout(0.1))
AlexNet.add(Dense(200, activation='relu'))
AlexNet.add(Dropout(0.1))
AlexNet.add(Dense(10,activation='softmax'))

AlexNet.compile(optimizer='SGD', loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
AlexNet.summary()

```



Model: "sequential\_4"

Layer (type)	Output Shape	Param #
=====		
conv2d_24 (Conv2D)	(None, 8, 8, 16)	448
max_pooling2d_12 (MaxPooling2D)	(None, 4, 4, 16)	0
conv2d_25 (Conv2D)	(None, 4, 4, 64)	2464
max_pooling2d_13 (MaxPooling2D)	(None, 2, 2, 64)	0
conv2d_26 (Conv2D)	(None, 2, 2, 64)	3264
conv2d_27 (Conv2D)	(None, 2, 2, 32)	1632
conv2d_28 (Conv2D)	(None, 2, 2, 20)	540
max_pooling2d_14 (MaxPooling2D)	(None, 1, 1, 20)	0
flatten_4 (Flatten)	(None, 20)	0
dense_8 (Dense)	(None, 200)	4200
dropout_16 (Dropout)	(None, 200)	0
dense_9 (Dense)	(None, 200)	40200
dropout_17 (Dropout)	(None, 200)	0
dense_10 (Dense)	(None, 10)	2010